

Errata

On Coppock, Leeper and Mullinix (2018), *Proceedings of the National Academy of Sciences* 115(49)

This note lists two corrections to “Generalizability of heterogeneous treatment effect estimates across samples,” published in the *Proceedings of the National Academy of Sciences* 115(49):12441–12446, doi [10.1073/pnas.1808083115](https://doi.org/10.1073/pnas.1808083115). Both are sentences that describe a column of a table incorrectly: one in the supplementary appendix, one in the main text. **Neither changes any conclusion of the paper**, and neither touches an estimate: every CATE, standard error, p-value and confidence limit in the article and in all 27 appendix tables reproduces exactly from the deposited archive.

The numbers in the corrected sentences and every number in the notes beneath them are computed from the deposited replication archive ([10.7910/DVN/4WNGEJ](https://doi.org/10.7910/DVN/4WNGEJ)) at the time this document is rendered. The code and output supporting the correction are at github.com/active-maintenance-aec/coppock_leeper_mullinix_2018.

1. The appendix’s Prop column is not the proportion the appendix says it is

The Prop column describes what proportion of the subjects in a given experiment belong to the associated covariate class.

Corrected:

The Prop column describes what proportion of the subjects in a given experiment belong to the associated covariate class, **divided by the number of terms in the regression that produced the row**.

The divisor is two in 23 of the 27 study pairs, whose subgroup regressions carry a treatment indicator and an intercept and nothing else, and three, four or five in the remaining four, whose regressions carry a further arm of the study’s own factorial design or the party identification its assignment probabilities depend on. An attribute’s classes therefore sum to between 0.20 and 0.50 across the tables rather than to 1.00. In appendix Table 1, for instance, the Original panel prints 0.18 for respondents aged 18 to 39, where they are in fact 0.35 of that sample.

The error is in the deposited code rather than in the typesetting: the expression that builds the column divides each subgroup's count by a total taken over every row of the fitted results rather than over the treatment coefficients alone, and all 787 printed values reproduce exactly under that expression. The correction restates what the column contains, because the alternative would be to reprint 27 tables in order to change one word in each cell. Nothing in the article's text, tables or figures quotes a Prop value, so no other published number depends on it.

2. Table 1 reports each study's sample size, not the number of observations the analyses use

The full list of studies, with the sample sizes used in the analyses reported here, is presented in Table 1.

Corrected:

The full list of studies, with **the number of respondents in each sample**, is presented in Table 1.

Table 1's columns count everyone in the survey wave a study was fielded in. The treatment indicator is defined only for the respondents that study assigned, so the analyses use fewer: 69,599 observations against the 101,745 the two columns total. The reported sample is the larger of the two in 18 of the 27 pairs and exactly equal in 8. It is smaller in 1, Johnston and Ballard (2016), whose outcome is measured over several items per respondent, so its observations outnumber its subjects. The gap is widest at Murtagh et al. (2012), where Table 1 reports 5,243 and the analyses use 640, or 12 per cent.

Several studies were fielded in the same wave, and a shared wave is counted once for each of them: the 27 MTurk entries take only 14 distinct values, four studies reporting the same 1,913 between them. Summing the column therefore adds a wave in as many times as it carried experiments, which is how 101,745 arises from 28,872 distinct MTurk wave sizes and 37,003 original-study respondents. That total is the figure the abstract and the significance statement give for the individual survey responses the pairs encompass. It is a faithful sum of Table 1 and reproduces exactly; readers should know that it counts wave memberships rather than distinct people, and that a respondent counted in a study's row was not necessarily assigned to that study's experiment.

The article contradicts itself on this point rather than merely being imprecise. For Bergan (2012), Table 1 reports 1,206 in the original sample, while every attribute in the corresponding panel of appendix Table 1 sums to 396. The appendix is right. No published estimate depends on the distinction: Table 1's counts are quoted in the article only as sample sizes, and each subgroup's own N is printed beside its estimate throughout the appendix.